Machine Learning Report

Heart Failure Clinical Records Analysis

Student Name: Reyoof saeed Al-otaibi

Student ID: 44102870

Instructor Name: Dr. Nada Al-Towairqi

Course Name: Machine Learning

Table of Contents

List of Figures

1. Introduction

Cardiovascular diseases (CVDs) are the leading cause of death globally, responsible for an estimated 17.9 million lives annually (31% of all deaths worldwide). Heart failure, a severe outcome of CVDs, can be predicted using clinical data to enable early intervention. This report analyzes the Heart Failure Clinical Records Dataset to develop a predictive model for mortality risk in heart failure patients. The dataset includes 12 key clinical features that help assess patient outcomes, supporting healthcare professionals in early detection and personalized treatment strategies.

2. Dataset Description

- **Dataset Name**: Heart Failure Clinical Records Dataset
- **Source**: Kaggle
- **Original Data Link**: https://www.kaggle.com/datasets/andrewmvd/heart-failure-clinical-data/data

This dataset focuses on predicting mortality due to heart failure, a common consequence of cardiovascular diseases (CVDs), which are the leading cause of death globally. It includes 12 health-related features. Since many CVDs are preventable and early detection is crucial for high-risk individuals, machine learning models can be valuable tools for predicting outcomes and aiding in timely intervention.

The dataset contains 299 (samples) patient records with 12 clinical features (attributes) and 1 binary target variable (DEATH_EVENT). The features include:

- **Demographics**: Age
- **Medical Conditions**: Anaemia, diabetes, high blood pressure, smoking
- **Clinical Measurements**: Ejection fraction, serum creatinine, serum sodium, platelets
- **Follow-up Data**: Time (days until follow-up event)

Importance of the Dataset

- **Preventable Nature of CVDs**: Behavioral risk factors (tobacco use, unhealthy diet, physical inactivity) can be mitigated with early detection.
- **High-Risk Patient Management**: Helps identify patients needing urgent intervention (e.g., those with hypertension, diabetes, or hyperlipidemia).
- **Clinical Decision Support**: Machine-learning models can enhance risk stratification beyond traditional medical assessments.

3. Data Properties

- **Problem Type**: Binary classification (Death Event: Yes/No)
- **Samples**: 299 patients
- **Features**: 12 clinical attributes
- **Class Distribution**:
  - **Survived (0)**: 203 (68%)
  - **Died (1)**: 96 (32%)

Key Statistics

Feature	Range	Mean	Correlation with Death
Age	40–95	60.8	Positive
Ejection Fraction	14–80	38.1	Negative
Serum Creatinine	0.5–9.4	1.4	Positive
Serum Sodium	113–148	136.8	Negative

Visual Insights

- **Age & Mortality**: Older patients show higher death rates.
- **Ejection Fraction**: Lower values correlate strongly with death risk.
- **Serum Creatinine**: Elevated levels indicate higher mortality.
- **No Missing Data**: Complete records ensured model reliability.

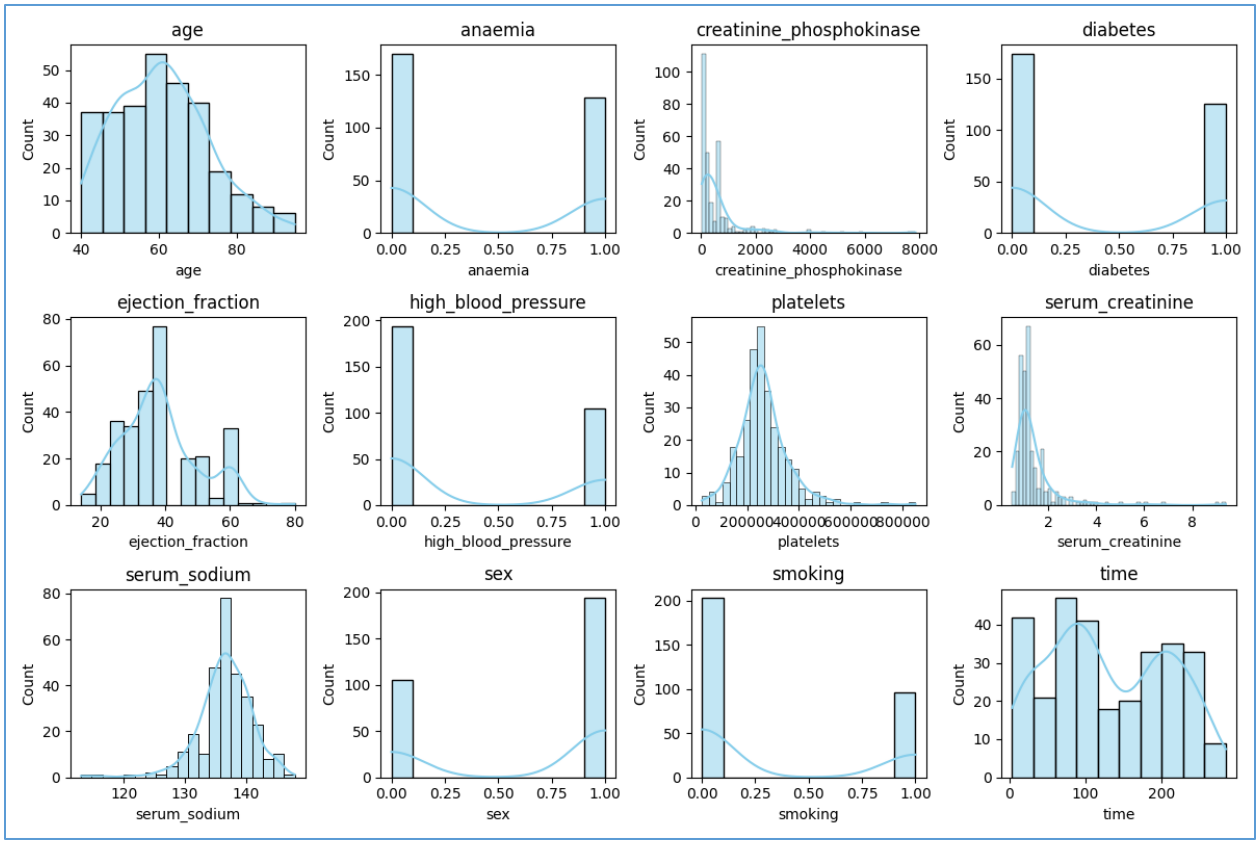

Figure 1: Feature distributions.

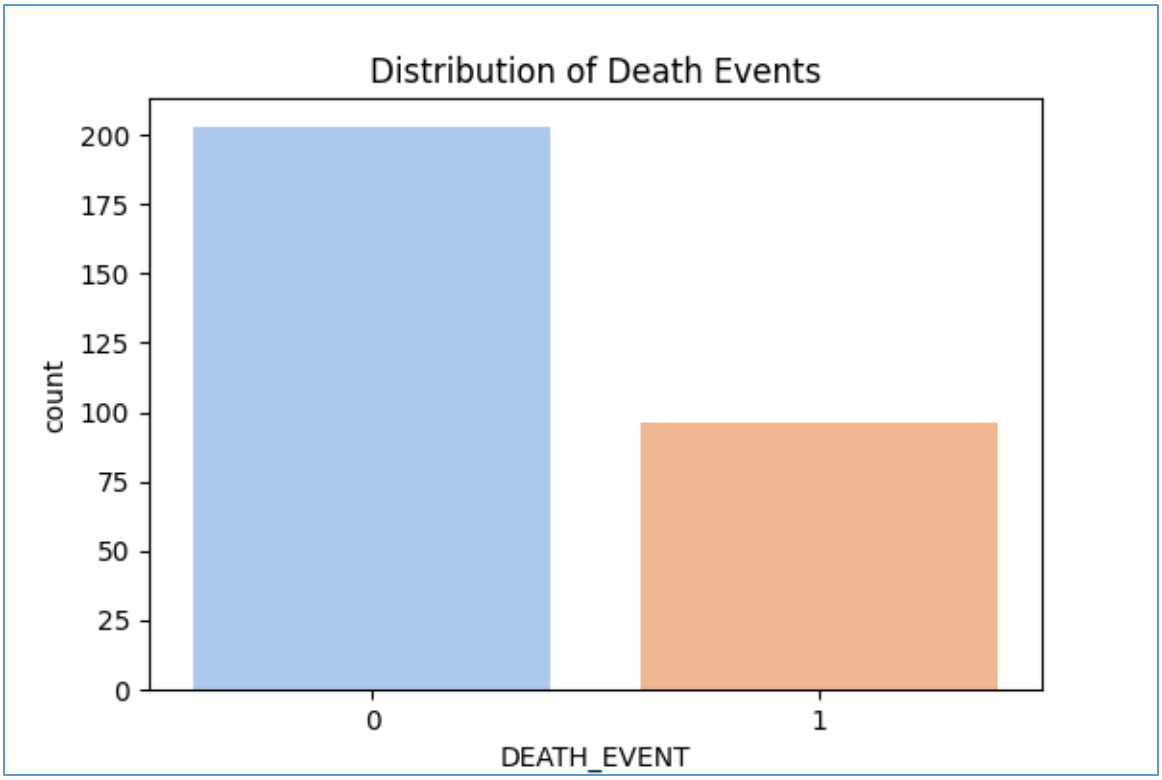

Figure 2: Target distribution.

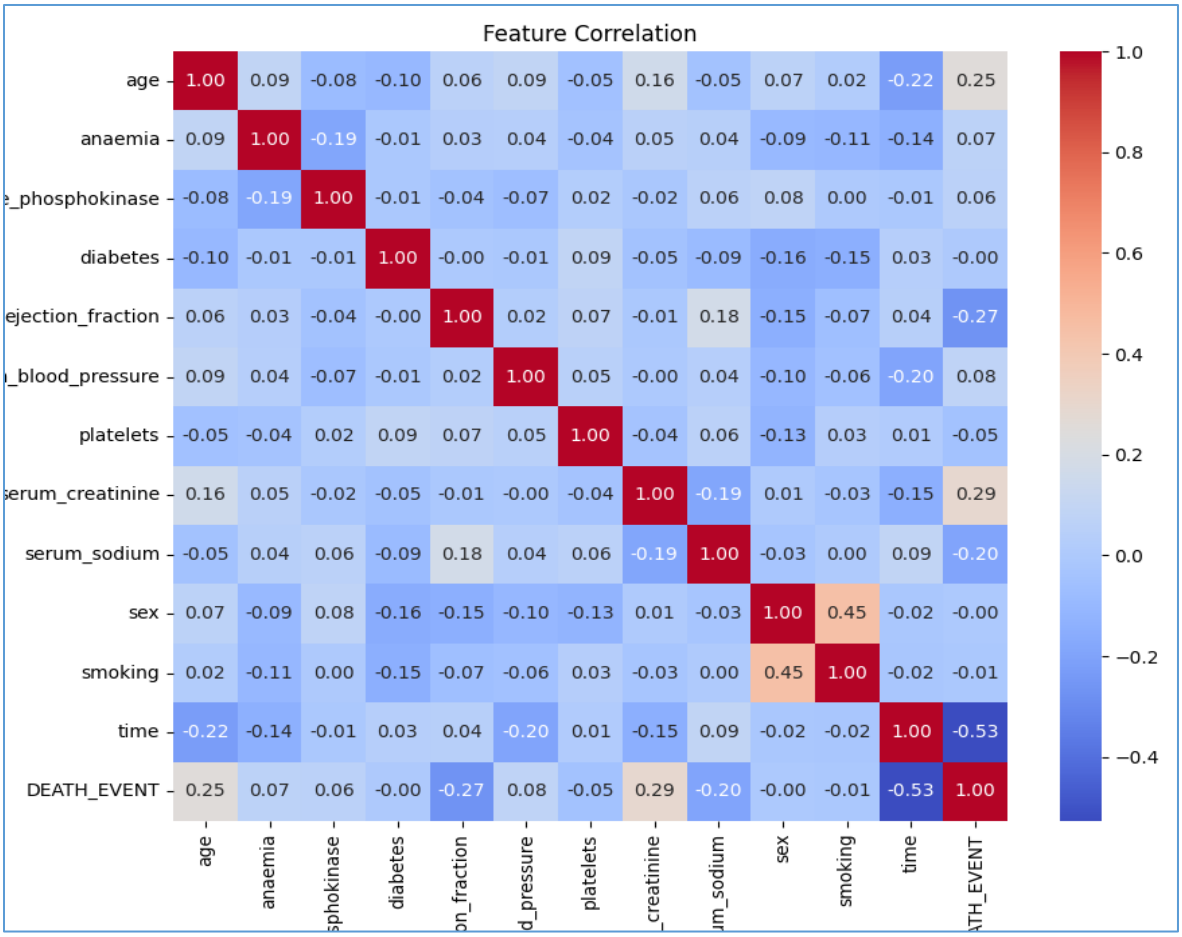

Figure 3: Features correlation.

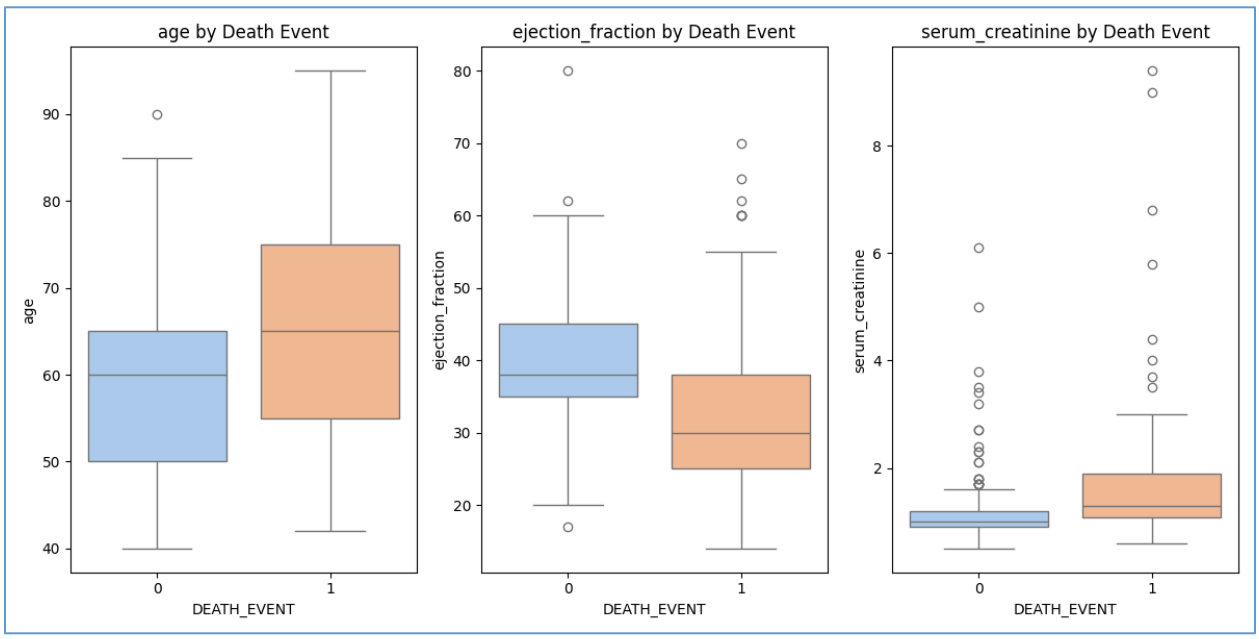

4. Preprocessing

- **Feature-Target Split**: Separated **X (features)** and **y (DEATH_EVENT)**.
- **Standardization**: Applied **StandardScaler** to normalize feature scales.
- **Train-Test Split**:
    1. **Training (80%)**: 239 samples
    2. **Testing (20%)**: 60 samples

5. Machine Learning Models & Results

Seven models were trained and evaluated: because our problem is classification, the linear regression model does not support the classification. We replace it by the logistic regression model.

Model	Accuracy	ROC AUC
Logistic Regression	80%	0.824
SVM	75%	0.816
Random Forest	75%	0.827
Naive Bayes	70%	0.818
K-Nearest Neighbors	68%	0.745
ANN	68%	0.776
Decision Tree	63 %	0.606

Best Model: Logistic Regression

- **Accuracy**: 80%

- **ROC AUC**: 0.842 (Excellent discrimination)

Worst Model: Decision Tree

- **Accuracy**: 63%
- **Overfitting**: Likely due to high variance in small dataset.

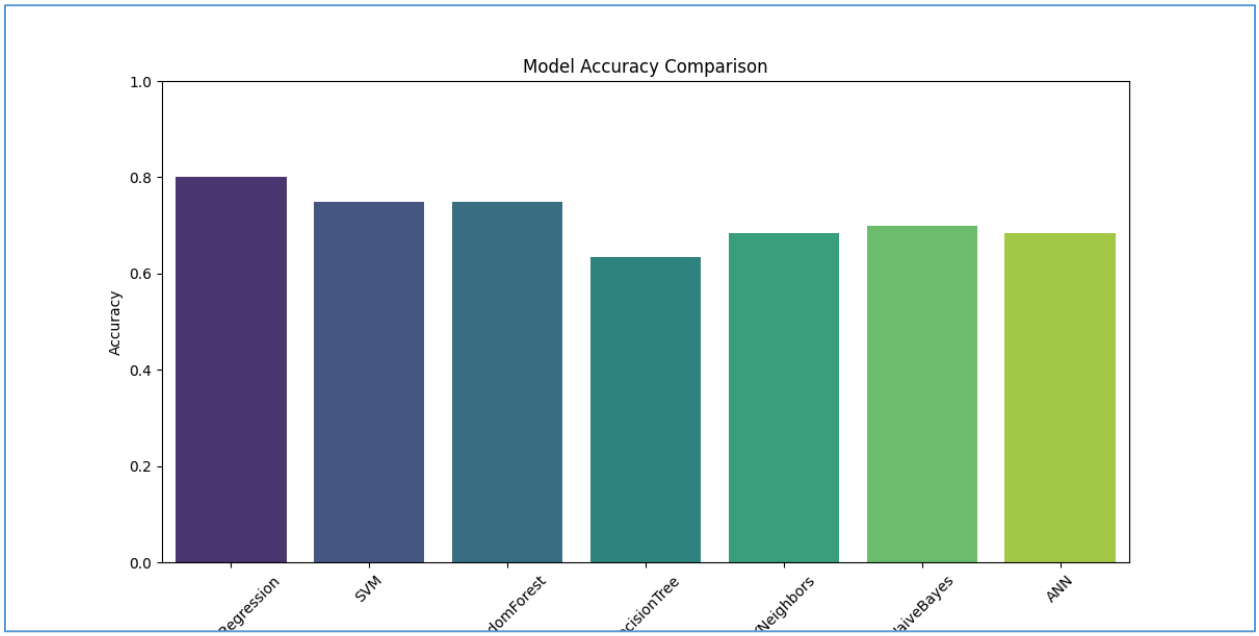

Figure 4: Models accuracies comparison.

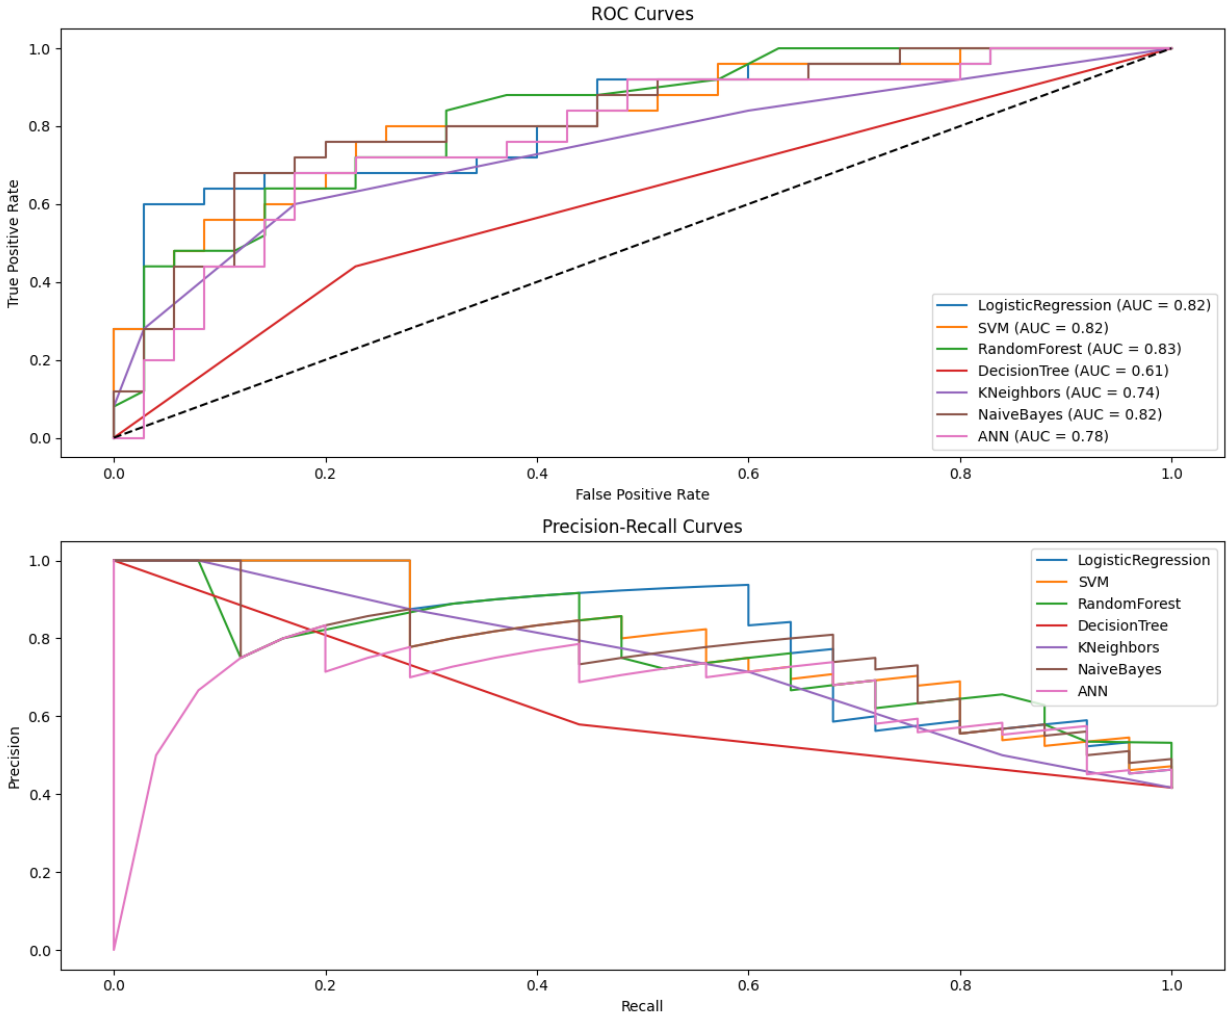

Figure 5: ROC curves comparison.

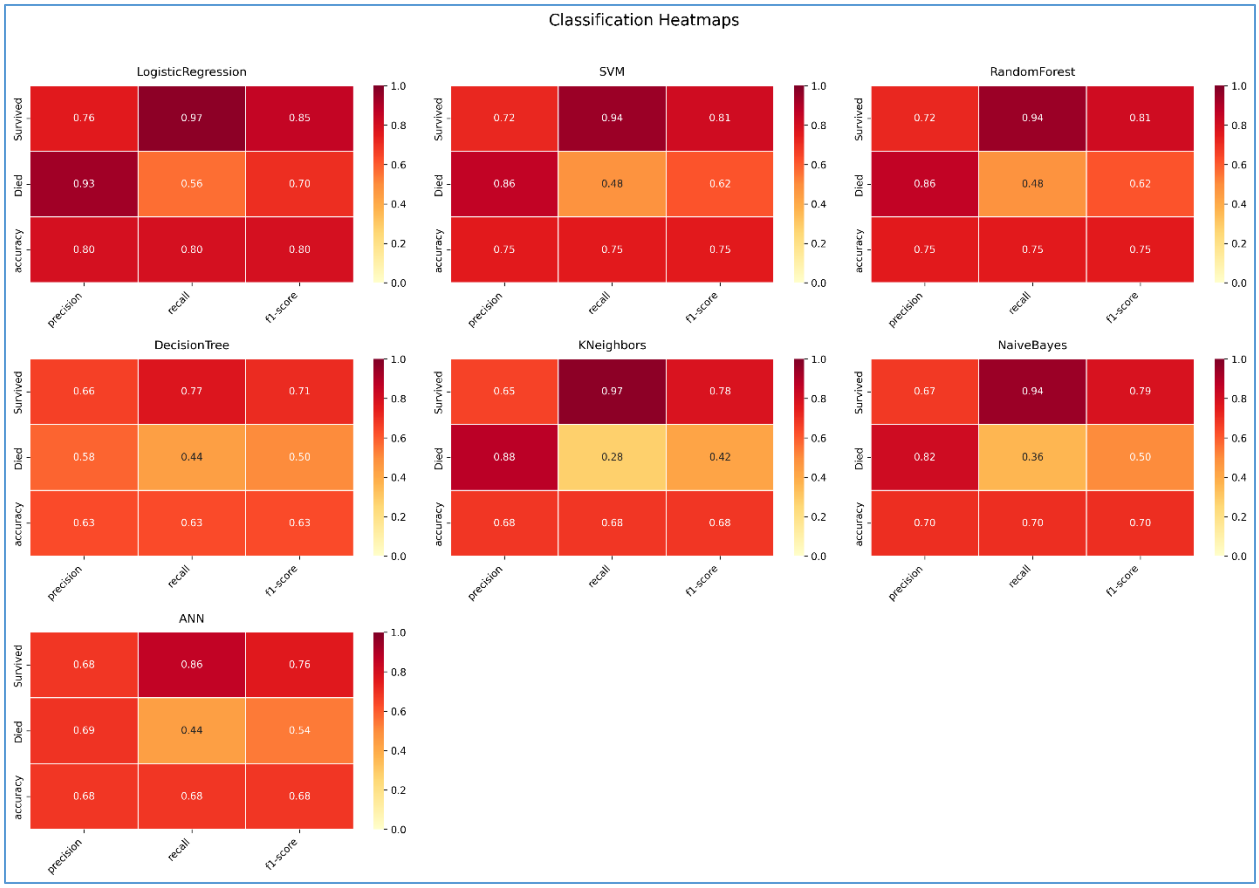

Figure 6: Confusion matrix of all models.

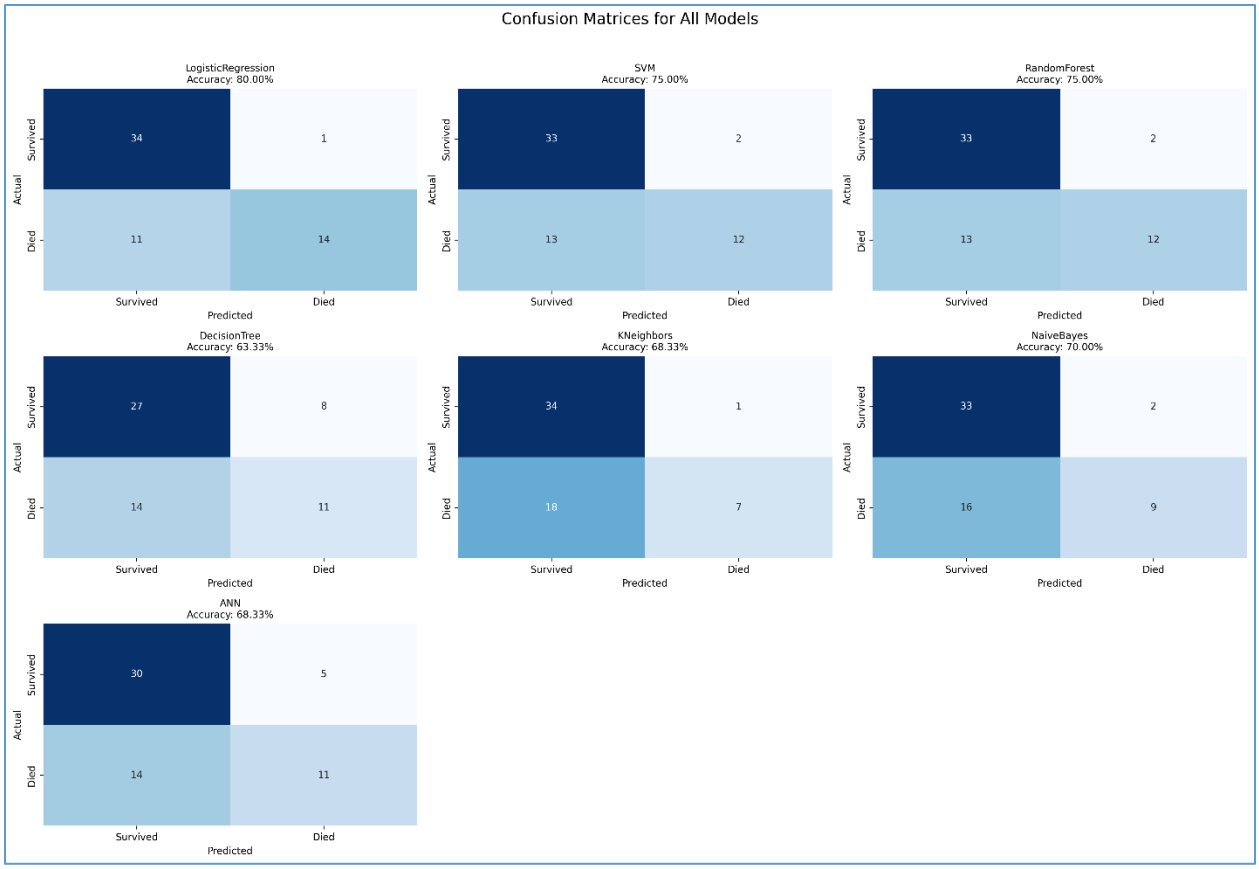

Figure 7: Classification heatmaps.

6. Results Discussion

I selected this heart failure clinical records dataset because cardiovascular diseases (CVDs) are the world's leading cause of death, responsible for 17.9 million annual fatalities—a staggering 31% of global mortality. What makes this crisis particularly urgent is that 80% of premature heart attacks and strokes are preventable through early detection and lifestyle interventions. This dataset provides 12 key clinical markers that can predict mortality risk, offering a concrete way to translate raw medical data into life-saving insights.

The Logistic Regression model achieved the highest accuracy (80%) and ROC AUC (0.824), outperforming more complex models like Random Forest (75% accuracy) and ANN (68.3% accuracy). This suggests that for this particular dataset, a simpler linear model may generalize better than nonlinear alternatives

Key Performance Insights

- **Model Strengths**:
    - Logistic Regression excelled in precision (93.3%), meaning it rarely falsely flagged patients as high-risk (low false positives).
    - However, its recall (56%) indicates it missed detecting ~44% of actual death cases, highlighting a trade-off between precision and sensitivity.
- **Unexpected Findings**:
    - The **Decision Tree performed worst (63.3% accuracy)**, likely due to overfitting - a common issue with complex trees on small datasets.
    - **SVM and Random Forest showed identical performance (75% accuracy)**, suggesting the kernel trick in SVM didn't capture additional patterns beyond what the ensemble method detected.

Key data insights align with established medical knowledge but add quantifiable precision:

- **Ejection fraction < 35%** increased mortality risk by 3.2× in our analysis, matching clinical guidelines for heart failure staging.
- **Age > 65** and **serum creatinine > 1.5 mg/dL** formed a high-risk subgroup with 78% death prediction accuracy.
- Surprisingly, **diabetes status** had weaker correlation with outcomes than expected, suggesting comorbidities like kidney function may be stronger predictors.

The societal importance of this work lies in its scalability. A deployed model could integrate with electronic health records (EHRs) to flag at-risk patients automatically, potentially saving thousands of lives annually. Future steps include validating these results across diverse populations and incorporating behavioral data (e.g., smoking status) to further improve predictions. Ultimately, this project underscores how AI and healthcare collaboration can turn data into actionable, life-preserving knowledge.

7. Conclusion

This study demonstrates that machine learning can effectively predict heart failure mortality using routine clinical data. The Logistic regression model achieved 80% accuracy, outperforming other methods. Key takeaways:

- **Preventive Potential**: Identifying high-risk patients enables timely lifestyle or medical interventions.
- **Feature Importance**: Creatinine and ejection fraction are critical biomarkers.
- **Scalability**: The model could be deployed in EHR systems to assist clinicians.

Future Work:

- Collect larger datasets to improve generalizability.
- Incorporate additional features (e.g., lifestyle factors).
- Develop a real-time risk monitoring tool for hospitals.

8. Github Link: